PRAMOD BURRA

Technical Lead | AI Engineer · Conversational AI · Agentic AI · LLM Engineering · BPM

Hyderabad, India · pramodburra@outlook.com · linkedin.com/in/pramodburra · pramodburra.com · github.com/pramod457

Technical Lead and AI Engineer with 11+ years, including 7 years on site in the US, designing, building and deploying production LLM powered applications, agentic and multi-agent workflows, and autonomous AI systems end to end. Strong in Python and Java with deep RAG, vector search, prompt engineering and LLM orchestration; ships enterprise grade services on Azure with Docker, Kubernetes and CI/CD, and builds automated LLM evaluation frameworks that track response quality, hallucination rate and decision accuracy. Currently architecting an AI native integration platform on Workato.

EXPERIENCE

Technical Lead / Kore.ai Developer, Nitco Inc Apr 2025 to present

Client: Enterprise conversational AI

- Architected and shipped enterprise-grade conversational AI for large-scale customer engagement and self-service: agent nodes, AI nodes, search AI, knowledge AI, dialog tasks and digital views, integrated with enterprise APIs, auth, CRM and backend microservices.
- Built RAG-style semantic search and knowledge retrieval with embeddings; engineered prompt strategies improving response quality, contextual grounding and hallucination reduction.
- Designed automated AI/LLM evaluation frameworks for intent accuracy, regression testing, and positive/negative test generation; optimized flows via confidence scoring, fallback analysis and containment metrics.
- Containerized and deployed bot and LLM services with Docker, CI/CD and observability. Logging plus response quality and hallucination monitoring keep production reliable and cost efficient.
- Currently designing an AI native integration platform on Workato: canonical process models, adapter based connectivity, and LLM assisted mapping and validation between source platforms and Workato recipes. The system architecture is the deliverable, including where AI is trusted to act and where it only proposes.

Java / Camunda Developer, Nitco Inc Mar 2023 to May 2025

Client: PepsiCo

- Built enterprise workflow-automation solutions on Camunda BPM and Java with scalable backend microservices and REST API integrations, reducing manual effort across business processes.

Java Developer, Nitco Inc Sep 2022 to Mar 2023

Client: BPX Energy

- Built automation and regression-testing frameworks in Java and Selenium, improving software quality through CI.

IBM ODM Consultant, Nitco Inc May 2015 to Mar 2023

Client: Humana · MetLife · PNC · Radian · NC Dept. of Revenue · PwC

- Designed enterprise business-rule and decision-automation services in IBM ODM, converting complex regulatory policies into scalable, highly available rule-based systems across healthcare, insurance, banking and government.
- Eight years of making system behaviour provable. That discipline now underpins the AI evaluation work.

SELECTED PROJECTS

AI-native iPaaS · An integration platform on Workato, built on Migration Studio

- The canonical model is the whole design. 53 entity types describe any integration process in neutral terms, so sources and targets never know about each other. Adding a platform is one adapter rather than a rewrite, and every Workato pattern built serves every source.
- Translation is rule based by choice. Rules are inspectable, repeatable and reviewable, which matters when the output is production integration code. This is exactly why the AI layer sits on top as assistance rather than underneath as the foundation.
- The target planner matches each canonical step to a Workato pattern before anything is generated, so coverage gaps are known up front instead of surfacing halfway through a build.
- Recipes are delivered stopped, never running. A generated integration that starts itself is a liability; a human turns it on after review.
- The revision store freezes every version with its findings, questions and evidence, so any generated step can be traced back to the exact source shape it came from.

- In design: where LLM assistance earns a place. Mapping suggestions and validation are the obvious candidates, because both produce something a person checks rather than something that silently ships.

DialogLens · Automated LLM evaluation & bot-testing platform

- The parser reads a bot definition into a conversation graph: dialog states as nodes, transitions as edges. Once the bot is a graph, test generation becomes traversal, which is why coverage scales with the bot instead of with the QA team.
- Three generators feed the suite. A deterministic walker covers every path; an FAQ generator covers knowledge answers that live outside the dialog tree; an entity permutation engine varies slot values so the same path is exercised with the inputs that actually break it.
- Execution is deliberately separated from generation. A pluggable executor drives whichever platform is under test (Kore.ai XO, Zendesk), so a suite written once survives a platform migration and the same conversations can be compared across vendors.
- Database first on Supabase and pgvector with zero local state, so any worker can pick up any run and results outlive the container. FastAPI services, containerized, CI driven.

ALO · Personal knowledge OS, a local first LLM application

- Ingestion transcribes, summarizes and extracts typed memories linked to entities: people, projects, companies, technologies.
- RAG over a knowledge graph with hybrid semantic + keyword search powers AI chat and briefs; content is reprocessed as better models ship.
- Local-first inference via Ollama with secure remote access.

SKILLS

Languages & APIs: Python, Java, SQL, JavaScript, REST API design, FastAPI, Spring Boot

LLM & GenAI: LLM orchestration, agentic AI, multi-agent workflows, prompt engineering, RAG, vector databases (pgvector, Supabase), semantic search, embeddings, knowledge graphs, function calling, guardrails, MCP, OpenAI and Azure OpenAI, LangChain, LangGraph

Conversational AI: Kore.ai XO platform, search AI, knowledge AI, agent and generative nodes, dialog tasks, intent and entity modeling, chatbots, autonomous workflows

Cloud & MLOps: Azure (OpenAI, AI Foundry, AI Search, AI Agent Service, Functions, App Services, DevOps), Docker, Kubernetes, CI/CD, GitHub Actions, Linux

Evaluation & Observability: Automated LLM evaluation, response quality and hallucination metrics, regression testing, monitoring, logging, performance and cost optimization

Enterprise: Microservices, event driven and distributed systems, Camunda BPM, IBM ODM, Workato, workflow automation, enterprise integrations

EDUCATION

Master of Science, Management Information Systems, Central Michigan University 2013 to 2015